



BINARY NUMBERS AND OTHER OBJECTS

Eric Prebys



“Numbers” in Computing

- As you'll learn if you take 118, everything in a computer ultimately reduces to 1s (TRUE) and 0s (FALSE)
 - These are internally defined by voltage levels.
- Numbers are made from groups individual signals, called “bits” (= “binary digit”)
- In the early days of computing, the exact format of numbers depended on the type of computer and there was a lot of variation:
 - Even “byte” was only defined as ”a group bits” when I was in college (sometimes 6, sometimes 8)



Types of Numbers

- In general, there are four types of numbers:
 - Unsigned Integers
 - Signed Integers
 - Fixed Point (really a special case of Signed Integers)
 - Floating Point
- Unsigned Integers are the only one with an unambiguous definition.
- Signed Integers and floating-point numbers have used different representations over the years.
- We'll focus on modern standards, but it's not impossible that you may encounter older standards somewhere down the road.



Unsigned integers (also used for addressing!)

- The binary representation of unsigned integers is simply the number expressed in base-2, in which the n^{th} bit represents 2^n (we'll always count bits starting at 0!)

0	0
1	1
10	2
11	3
100	4
101	5
110	6
111	7

1000	8
1001	9
1010	10
1011	11
1100	12
1101	13
1110	14
1111	15

- A n -bit word has a range 0 to $(2^n - 1)$



Converting from Binary to Decimal

- Simply add up the values of the individual digits

$$\begin{array}{r} 128 \ 64 \ 32 \ 16 \ 8 \ 4 \ 2 \ 1 \\ | \quad | \\ 1 \ 0 \ 0 \ 1 \ 1 \ 0 \ 1 \ 1 \\ \hline 128 + 0 + 0 + 16 + 8 + 0 + 2 + 1 \\ = 155 \end{array}$$



Converting from Decimal to Binary

- Recursive subtraction
 - Find the largest value of n for which 2^n is less than the number
 - Set bit n
 - Subtract 2^n from original number
 - Repeat until the remainder is zero or $n=0$ (ie, remainder=1)
- Using our previous example (155)
 - $2^7=128$ is the largest power of 2 smaller than 155. Set bit 7. $155-128=27$
 - $2^4=16$ is the largest power of 2 smaller than 27. Set bit 4. $27-16=11$
 - $2^3=8$. Set bit 3. $11-8=3$
 - $2^1=2$. Set bit 1. $3-2=1$
 - $2^0=1$. Set bit 0

$$\begin{array}{cccccccc} 128 & 64 & 32 & 16 & 8 & 4 & 2 & 1 \\ \downarrow & \downarrow \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 & 1 \end{array}$$

$$128 + 0 + 0 + 16 + 8 + 0 + 2 + 1 = 155$$

- Or just ask Google!



Signed Integers

- The three most common ways of representing signed integers

- Sign and magnitude

- Very old computers
- The mantissa of floating point numbers

- Offset binary

$$N' = N + (2^{n-1} - 1)$$

- Used for some hardware
- Used for exponent in floating point

- Twos Complement

$$-N = 2^n - N$$

- Standard in modern computers
- Addition is always the same

SIGN AND MAGNITUDE		OFFSET BINARY		TWO'S COMPLEMENT	
Decimal	Bit Pattern	Decimal	Bit Pattern	Decimal	Bit Pattern
7	0111	8	1111	7	0111
6	0110	7	1110	6	0110
5	0101	6	1101	5	0101
4	0100	5	1100	4	0100
3	0011	4	1011	3	0011
2	0010	3	1010	2	0010
1	0001	2	1001	1	0001
0	0000	1	1000	0	0000
0	1000	0	0111	-1	1111
-1	1001	-1	0110	-2	1110
-2	1010	-2	0101	-3	1101
-3	1011	-3	0100	-4	1100
-4	1100	-4	0011	-5	1011
-5	1101	-5	0010	-6	1010
-6	1110	-6	0001	-7	1001
-7	1111	-7	0000	-8	1000
		16 bit range -32,767 to 32,767		16 bit range -32,767 to 32,768	
					16 bit range -32,768 to 32,767



Binary addition

- Binary addition is just like decimal addition except that you carry a whole lot more. 8-bit example

$$\begin{array}{r} \text{ } & \nwarrow & \nwarrow \nwarrow \\ 00100011 & & \\ + & 10110111 & \\ \hline 11011010 & & \end{array}$$

- Addition can "overflow", in which case, higher order bit(s) will be ignored

$$\begin{array}{r} 10100011 \\ + 10110111 \\ \hline \cancel{X}01011010 \end{array}$$

This would be dropped for 8-bit words



Two's Complement

- An n -bit negative number is formed by subtracting the number from 2^n
- In additions, overflow bits are ignored
- Example: $-55 (-110111_2)$ in 8 bits

$$2^8 - 55 = 256 - 55 = 201 = 11001001_2$$

- Trick: just invert all bits (“1’s complement”) and add 1!
- Check

To extend to higher number of bits, repeat MSB
In 16 bits

$$\begin{array}{r} 00110111_2 \\ + \quad 11001001_2 \\ \hline \text{X}00000000_2 \end{array}$$

$$\begin{aligned} 55 &= 0000000000110111_2 \\ -55 &= 111111111001001_2 \end{aligned}$$



Binary Multiplication

- Binary multiplication is just like decimal multiplication, except that multiplying something by 2^n just shifts it n bits to the left.

$$1011011_2 \cdot 100_2 = 101101100_2$$

1101 (13)
×1011 (11)
—————
1101
1101
1101
1101

1101
+11010
—————
100111
+110100
—————
10001111 (143)

- Note! Just like decimal multiplication, multiplying two n-bit numbers can result in a 2n-bit number -> easy to overflow



Fixed Point Numbers

- Fixed point numbers are just an offset case of signed integers

128	64	32	16	8	4	2	1			0.5	0.25	0.125	0.0625	0.03125	0.015625	0.0078125	0.00390625
128	64	32	16	8	4	2	1			$\frac{1}{2}$	$\frac{1}{4}$	$\frac{1}{8}$	$\frac{1}{16}$	$\frac{1}{32}$	$\frac{1}{64}$	$\frac{1}{128}$	$\frac{1}{256}$
0	0	1	0	1	0	1	1			0	1	0	0	1	0	1	1

- Example: 37.55

- Get roughly equivalent precision by expressing in 128ths (7 bits), but it's common to work in groups of 4 (16ths) or 8 (256ths), so let's use 8 bits after the radix:

- Multiply by 2⁸ (256) $37.55 \times 256 = 9613$

- Convert to binary $9613_{10} = 10010110001101_2$

- Shift radix left by 8

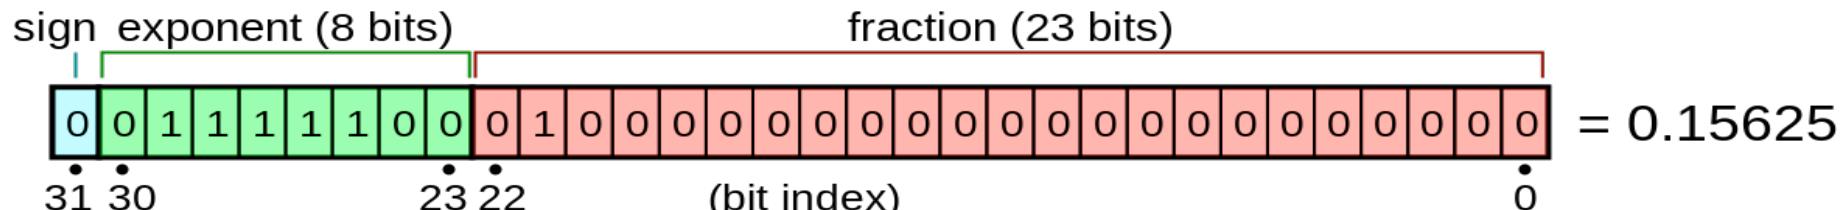


$$\frac{37}{256} \quad \frac{141}{256} = .551$$



Floating Point Numbers (IEEE 754*)

- Single precision (32-bit)



$$(-1)^{b_{31}} \times (1.b_{22}b_{21}\dots b_0)_2 \times 2^{(b_{30}b_{29}\dots b_{23})_2 - 127}$$

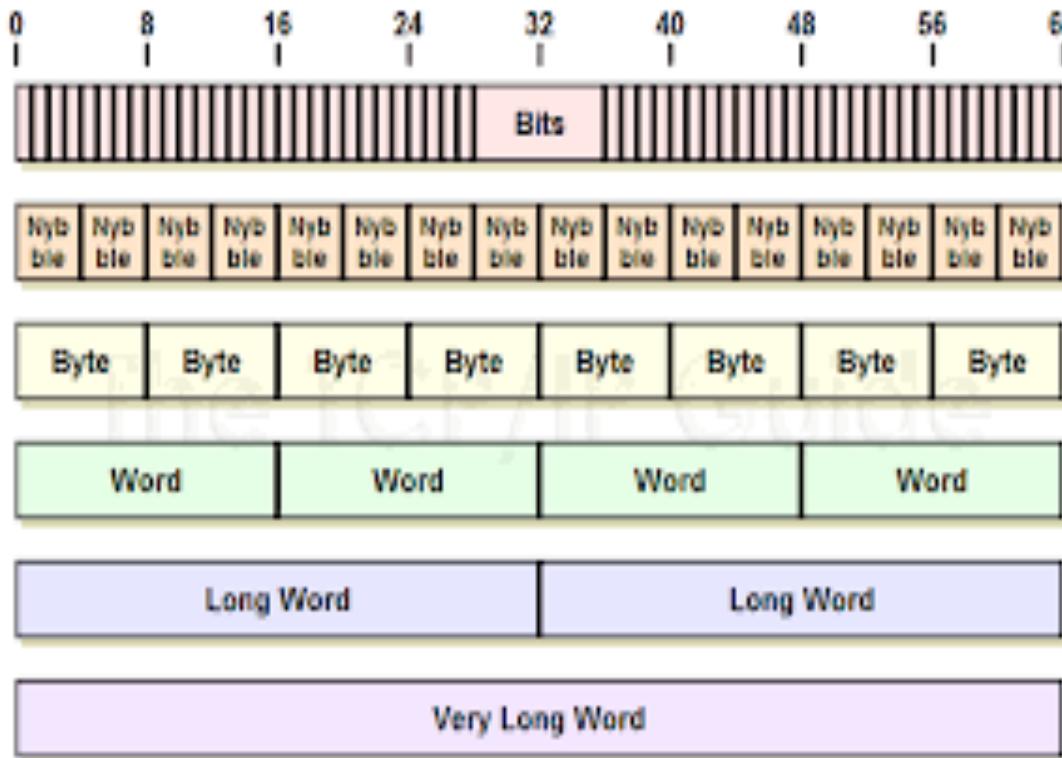
Leading 1 assumed (why?)

- IEEE 754 includes standards from 16-bit (half precision) to 256-bit (octuple precision) floating point.
- Special values
 - exponent bits set to zero \rightarrow zero (also allows for -0!)
 - Exponent bits all set and all fraction bits clear \rightarrow Infinity
 - Exponent bits all set and *some* fraction bits set \rightarrow NaN

*became a standard in 1985



Grouping Bits in Modern Computers



This took a long time to evolve. Examples

- Some old computers (IBM, PDP) had 6-bit bytes
- CDC mainframes had 60-bit words, with totally different encoding.

- Hardware will generally not group bits exactly like this
 - Still tend to group things in 8-bit bytes
 - The term “word” is used generically



Representing Groups of Bits

- Regardless of their use, groups of bits can always be compactly represented with hexadecimal numbers
 - Each byte is represented by a two-digit hex number 00->FF, where each digit represents 4 bits

Digit	Bits	Digit	Bits
0	0000	8	1000
1	0001	9	1001
2	0010	A	1010
3	0011	B	1011
4	0100	C	1100
5	0101	D	1101
6	0110	E	1110
7	0111	F	1111

- We can then represent larger groups of bits by larger hex numbers, while still respecting byte boundaries
 - One byte is two hex digits
 - Historically, 6-bit bytes were represented by two digit octal.



Representing Characters

- The most common character encoding is American Standard Code for Information Interchange (ASCII), in which each 8-bit byte represents one character
- In recent years, the 16-bit “Unicode” standard has emerged to accommodate foreign alphabets as well as a wide range of special characters.

USASCII code chart																
b ₇	b ₆	b ₅	b ₄	b ₃	b ₂	b ₁	b ₀	Column Row	0	1	2	3	4	5	6	7
0	0	0	0	0	0	0	0	NUL	DLE	SP	0	@	P	'	p	
0	0	0	0	1	1	1	1	SOH	DC1	!	1	A	Q	a	q	
0	0	1	0	2	2	2	2	STX	DC2	"	2	B	R	b	r	
0	0	1	1	3	3	3	3	ETX	DC3	#	3	C	S	c	s	
0	1	0	0	4	4	4	4	EOT	DC4	\$	4	D	T	d	t	
0	1	0	1	5	5	5	5	ENQ	NAK	%	5	E	U	e	u	
0	1	1	0	6	6	6	6	ACK	SYN	B	6	F	V	f	v	
0	1	1	1	7	7	7	7	BEL	ETB	'	7	G	W	g	w	
1	0	0	0	8	BS	CAN	(8	H	X	h	x				
1	0	0	1	9	HT	EM)	9	I	Y	i	y				
1	0	1	0	10	LF	SUB	*	:	J	Z	j	z				
1	0	1	1	11	VT	ESC	+	:	K	[k	{				
1	1	0	0	12	FF	FS	,	<	L	\	l	l				
1	1	0	1	13	CR	GS	-	=	M]	m)				
1	1	1	0	14	SO	RS	.	>	N	^	n	~				
1	1	1	1	15	SI	US	/	?	0	—	o	DEL				



Evolution of Precision

- Numbers used to default to 16-bit precision
 - Unsigned range: 0 to 65,535
 - Signed range: -32,768 to +32,767
 - This was standard for PC until the mid-1990s
- All systems eventually moved to 32-bit
 - Unsigned range: 0 to ~4 billion
 - Signed range: -2 billion to +2 billion
- All modern systems are 64-bit
 - Unsigned range: 0 to 18 quintillion
 - Signed range: -9 quintillion to 9 quintillion
- The current release of Python is 64-bit, but Python handles integers in a special way that allows arbitrarily large range.



Integer Mathematical Precision

- Integer math is exact!
 - You'll always get the correct answer until it overflows.
 - For signed 64-bit this is a pretty large number, but we deal with very large numbers.
 - Can happen pretty quickly with multiplication.
- Python “int” objects have arbitrary size.
 - Necessary size is determined when the object is initialized
 - Numerous internal methods give details about the structure
 - Integer operations will *never* overflow or have round-off errors!
 - (Jupyter example shortly)



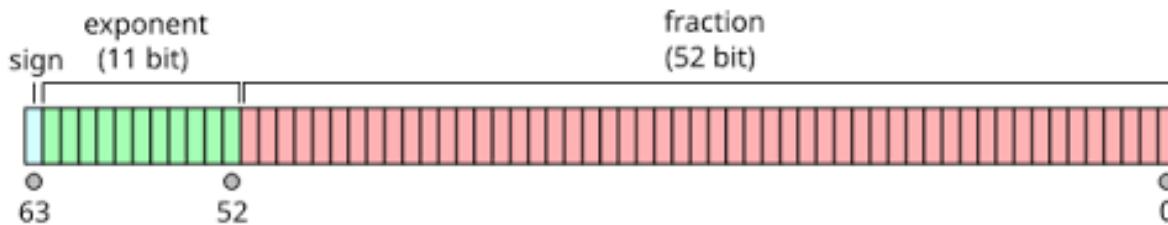
Important differences between Python and Other OO Languages...

- Both C++ and Java distinguish between “primitive” data types and “objects”
 - Primitives: integer, floating point, character, etc. These have no inherent methods!
 - Objects: More complex structures that include data and associated methods.
- Python has *only* objects, even for simple data types:
 - Everything has methods.
 - This leads to more versatility
 - The distinction is between “mutable” and “immutable”, which we’ll discuss shortly.



Floating Point Precision in Python

- While C/C++ and Java distinguish between "float" (32-bit) and "double" (64-bit) numbers, 64-bit Python "float" is *always* 64-bit.



- Range: $\pm 4.9 \times 10^{-324}$ to $\pm 1.8 \times 10^{308}$
- Precision: 52 binary digits \sim 15 decimal digits
- Multiplication
 - Add exponents
 - Multiply fractions (including implicit 1)
 - Still only have 52-bit precision
- Addition
 - Right-shift smaller number so exponent matches larger number
 - Add fractions (including implicit 1)
 - Will round off if one number is $> \sim 2^{53}$ smaller than the other.



Other Number Types

- When it comes to built-in number types, Python has ONLY
 - **int**: unlimited size integers (math *extremely* inefficient)
 - **float**: floating point with precision of the OS (now always 64-bit)
- Sometimes, we'll want other things
 - Integers of fixed bit length
 - Floating point numbers with different bit lengths (32 or 128)
 - This is *extremely* important when calling routines in C/C++ and other languages.
- This functionality will be provided by the “Numerical Python” (numpy) library, which we’ll introduce shortly.
- Boolean (**bool**) variable
 - Values of “True” or “False”. Can be thought of as a 1-bit integer



tuples, lists, and keyed lists (dicts)

- Ordered sets:
 - “tuples” are ordered sets of objects in which individual items CANNOT be modified
 - “lists” are ordered sets of objects in which individual items.
 - Lists of numbers are referred to as “arrays”, but we’ll do our array math with numerical python (numpy)
- Keyed sets
 - Sets in which elements are addressed by keyword are called “keyed lists” or “dicts”



Mutable vs. Immutable

- Simple objects in Python are “immutable”, meaning they can’t be changed once created.
 - This is generally associated with “pass by value”, although the way Python implements it is kind of weird.
- More complex objects are “mutable”, meaning that objects can be changed once they are created
 - When one object is equated to another, it is “passed by reference” rather than creating a new object.

(go to notebook demos)